

S3 - KD-TREES

Assignment grading.

- Correctness of `PointSET.java` (Mooshak score): 15%
- Correctness of `KdTree.java` (Mooshak score): 50%
- Efficiency of the `KdTree.java` implementation: 20%
- Report: 15%
- (Bonus) Implementation of `draw()` 5%

Handin Test `PointSET.java` and `KdTree.java` on Mooshak. *Note:* We will supply `RectHV.java`, `Point2d.java`, and `algs4.jar`. Your programs may not call any library functions outside `algs4.jar` other than those permitted in “Teaching Methods” on Canvas. Submit on Canvas:

1. The final `.java` files (`PointSET.java` and `KdTree.java`).
2. The report (in PDF, or in `readme.txt`).
3. (Bonus) Visual output when run on `circle10.txt` and `input10K.txt`

This assignment is for groups of up to 2 persons.

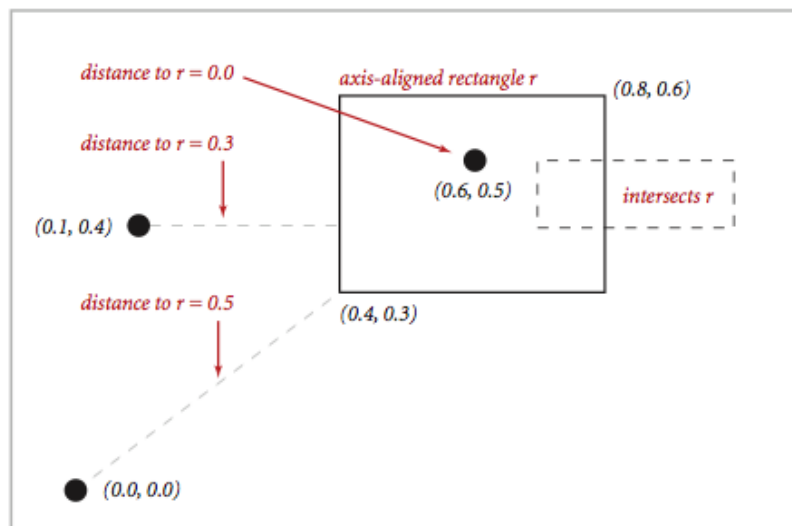
BRUTE-FORCE – POINTSET

Write a data type to represent a set of points in the unit square (all points have x- and y-coordinates between 0 and 1) using a 2d-tree to support efficient range search (find all of the points contained in a query rectangle) and nearest neighbor search (find a closest point to a query point). 2d-trees have numerous applications, ranging from classifying astronomical objects to computer animation to speeding up neural networks to mining data to image retrieval.

Geometric primitives. To get started, use the following geometric primitives for points and axis-aligned rectangles in the plane.



MYND 1.



Use the immutable data type `Point2D.java` (part of `algs4.jar`) for points in the plane. Here is the subset of its API that you may use:

```

1 public class Point2D implements Comparable<Point2D> {
2     // construct the point (x, y)
3     public Point2D(double x, double y);
4     public double x(); // x-coordinate
5     public double y(); // y-coordinate
6     // Euclidean distance between two points
7     public double distanceTo(Point2D that);
8     // square of Euclidean distance between two points
9     public double distanceSquaredTo(Point2D that);
10    // for use in an ordered symbol table
11    public int compareTo(Point2D that);
12    // does this point equal that?
13    public boolean equals(Object that);
14    // draw to standard draw
15    public void draw();
16    // string representation
17    public String toString();
18 }

```

Use the immutable data type RectHV.java (not part of algs4.jar) for axis-aligned rectangles. Here is the subset of its API that you may use:

```

1 public class RectHV {
2     // construct the rectangle [xmin, xmax] x [ymin, ymax]
3     // throw a java.lang.IllegalArgumentException
4     // if (xmin > xmax) or (ymin > ymax)
5     public RectHV(double xmin, double ymin, double xmax, double ymax);
6     // minimum x-coordinate of rectangle
7     public double xmin();
8     // minimum y-coordinate of rectangle
9     public double ymin();
10    // maximum x-coordinate of rectangle
11    public double xmax();
12    // maximum y-coordinate of rectangle
13    public double ymax();
14    // does this rectangle contain the
15    // point p (either inside or on boundary)?
16    public boolean contains(Point2D p);
17    // does this rectangle intersect
18    // that rectangle (at one or more points)?
19    public boolean intersects(RectHV that);
20    // Euclidean distance from
21    // point p to the closest point in rectangle
22    public double distanceTo(Point2D p);
23    // square of Euclidean distance from
24    // point p to closest point in rectangle
25    public double distanceSquaredTo(Point2D p);
26    // does this rectangle equal that?
27    public boolean equals(Object that);
28    // draw to standard draw
29    public void draw();
30    // string representation
31    public String toString();
32 }

```

Do not modify these data types.

Assignment Write a mutable data type PointSET.java that represents a set of points in the unit square. Implement the following API by using a red-black BST (using either SET from algs4.jar or java.util.TreeSet).

```

1 public class PointSET {
2     // construct an empty set of points
3     public PointSET()
4
5     // is the set empty?
6     public boolean isEmpty()
7
8     // number of points in the set
9     public int size()
10
11    // add the point p to the set (if it is not already in the set)
12    public void insert(Point2D p)
13
14    // does the set contain the point p?
15    public boolean contains(Point2D p)
16
17    // draw all of the points to standard draw
18    public void draw()
19
20    // all points in the set that are inside the rectangle
21    public Iterable<Point2D> range(RectHV rect)
22

```

```

23 // a nearest neighbor in the set to p; null if set is empty
24 public Point2D nearest(Point2D p)
25 }


```

Your implementation should support `insert()` and `contains()` in time proportional to the logarithm of the number of points in the set in the worst case; it should support `nearest()` and `range()` in time proportional to the number of points in the set.

2D-TREE IMPLEMENTATION

Write a mutable data type `KdTree.java` that uses a 2d-tree to implement the same API (but replace `PointSET` with `KdTree`). A 2d-tree is a generalization of a BST to two-dimensional keys. The idea is to build a BST with points in the nodes, using the x- and y-coordinates of the points as keys in strictly alternating sequence.

- *Search and insert.* The algorithms for search and insert are similar to those for BSTs, but at the root we use the x-coordinate (if the point to be inserted has a smaller x-coordinate than the point at the root, go left; otherwise go right); then at the next level, we use the y-coordinate (if the point to be inserted has a smaller y-coordinate than the point in the node, go left; otherwise go right); then at the next level the x-coordinate, and so forth.



MYND 2.

- *Draw.* A 2d-tree divides the unit square in a simple way: all the points to the left of the root go in the left subtree; all those to the right go in the right subtree; and so forth, recursively. Your `draw()` method should draw all of the points to standard draw in black and the subdivisions in red (for vertical splits) and blue (for horizontal splits). This method need not be efficient—it is primarily for debugging.

The prime advantage of a 2d-tree over a BST is that it supports efficient implementation of range search and nearest neighbor search. Each node corresponds to an axis-aligned rectangle in the unit square, which encloses all of the points in its subtree. The root corresponds to the unit square; the left and right children of the root corresponds to the two rectangles split by the x-coordinate of the point at the root; and so forth.

- *Range search.* To find all points contained in a given query rectangle, start at the root and recursively search for points in both subtrees using the following pruning rule: if the query rectangle does not

intersect the rectangle corresponding to a node, there is no need to explore that node (or its subtrees). A subtree is searched only if it might contain a point contained in the query rectangle.

- *Nearest neighbor search.* To find a closest point to a given query point, start at the root and recursively search in both subtrees using the following pruning rule: if the closest point discovered so far is closer than the distance between the query point and the rectangle corresponding to a node, there is no need to explore that node (or its subtrees). That is, a node is searched only if it might contain a point that is closer than the best one found so far. The effectiveness of the pruning rule depends on quickly finding a nearby point. To do this, organize your recursive method so that when there are two possible subtrees to go down, you always choose the subtree that is on the same side of the splitting line as the query point as the first subtree to explore—the closest point found while exploring the first subtree may enable pruning of the second subtree.

README AND OTHER USEFUL INFORMATION

You may use the following interactive client programs to test and debug your code.

- `KdTreeVisualizer.java` computes and draws the 2d-tree that results from the sequence of points clicked by the user in the standard drawing window.
- `RangeSearchVisualizer.java` reads a sequence of points from a file (specified as a command-line argument) and inserts those points into a 2d-tree. Then, it performs range searches on the axis-aligned rectangles dragged by the user in the standard drawing window.
- `NearestNeighborVisualizer.java` reads a sequence of points from a file (specified as a command-line argument) and inserts those points into a 2d-tree. Then, it performs nearest neighbor queries on the point corresponding to the location of the mouse in the standard drawing window.

Address the following issues in your report:

- Give the total memory usage in bytes (using tilde notation) of your 2d-tree data structure as a function of the number of points N , using the memory-cost model from lecture and Section 1.4 of the textbook. Count all memory that is used by your 2d-tree, including memory for the nodes, points, and rectangles.
- Give the expected running time in seconds (using tilde notation) to build a 2d-tree on N random points in the unit square. (Do not count the time to read in the points from standard input.)
- How many nearest neighbor calculations can your 2d-tree implementation perform per second for `input100K.txt` (100,000 points) and `input1M.txt` (1 million points), where the query points are random points in the unit square? (Do not count the time to read in the points or to build the 2d-tree.)
- Repeat the last question with the brute-force implementation.

This assignment was developed by Kevin Wayne.

SCHOOL OF COMPUTER SCIENCE, REYKJAVÍK UNIVERSITY, MENNTAVEGI 1, 101 REYKJAVÍK

E-mail address: `mmh@ru.is`